

Testing Resiliency of AWS Services using the Concept of Chaos Engineering

Dr. Purushottam J Assudani, Pranjal Ostwal, Pranav Kumbhalkar and Rajat Kale
Shri Ramdeobaba College of Engineering and Management Nagpur, Maharashtra, India

Abstract— Ensuring software availability in the cloud is a critical aspect of software development across diverse domains. Uninterrupted website performance is a fundamental requirement for successful development organizations. During peak periods, such as high-traffic sales events, server failures can lead to customer dissatisfaction. Platform-as-a-Service (PaaS) providers offer solutions to scale websites and improve uptime, necessitating comprehensive testing.

This paper aims to review different ways of performing chaos engineering in AWS environment like: use of third party software like Gremlin, a service provided by AWS i.e. Fault Injection Simulator (FIS), use of software development kits like python SDK in various services provided by AWS, including Elastic Compute Cloud (EC2), Relational Database Services (RDS), and Availability Zones. By addressing challenges related to service availability, this research contributes to enhancing the reliability and performance of cloud-based applications, ensuring seamless user experiences and minimizing potential disruptions.

Keywords— *Cloud-based software availability, Amazon Web Services (AWS), Elastic Compute Cloud (EC2), Relational Database Services (RDS), Availability Zone (AZ), PaaS, automated script, failure simulation, software testing, cloud computing, Simple State Machine (SSM).*

I. INTRODUCTION

Chaos Engineering, a discipline focused on experimenting with systems to build confidence in their resilience, has gained significant attention in the world of cloud computing. Cloud computing, an internet-based distributed computing technology, offers a new paradigm for dynamically provisioning computing services. With its dynamic scalability and utilization of virtualized resources, cloud computing has become an indispensable technique for many organizations. However, optimizing energy efficiency in the cloud while maintaining quality of service (QoS) is a challenge that requires a trade-off between performance and energy consumption.

In the realm of cloud computing, chaos engineering serves as a resilience guardian, allowing organizations to simulate controlled disruptions and observe the behaviour of their systems in a safe environment. By injecting real-world disruptions into various components of the cloud infrastructure, such as at the cloud-provider, infrastructure, workload-component, and process levels, teams can learn from faults and enhance the resilience of their workloads. Chaos engineering helps validate the effectiveness of observability tooling, incident response,

and recovery mechanisms, ensuring minimal customer impact.

Under the AWS Well-Architected framework, which provides guidance for designing and implementing scalable cloud architectures, the reliability pillar emphasizes the importance of chaos engineering. It highlights distributed system design, recovery planning, and adaptability to changing requirements as crucial factors for building reliable cloud systems. Additionally, AWS offers the AWS Fault Injection Simulator (AWS FIS), a fully managed service that enables running fault injection experiments to simulate real-world AWS faults. By integrating chaos engineering principles into their software development lifecycle and CI/CD pipelines, organizations can continuously enhance their systems' resilience, validate logs and metrics, and ensure effective incident response and recovery capabilities.

While cloud computing revolutionizes the delivery of hosted services over the internet, it is the combination of chaos engineering and the AWS Well-Architected framework that empowers organizations to build secure, high-performing, and resilient cloud infrastructures. By proactively identifying and mitigating potential disruptions, improving workload resilience, and validating monitoring mechanisms, businesses can confidently navigate the complexities of the cloud and ensure optimal system performance, security, energy efficiency, and cost-effectiveness.

II. LITERATURE SURVEY

The emerging field we are discussing has gained significant importance recently, despite being relatively new. The research conducted in this area has had a profound impact. Even with advanced computing power, it becomes futile if end-users cannot fully utilize website features due to failures. In this context, we will explore some recent research works and their practical applications.

Let's start by examining a significant paper published in 2019 that serves as a key reference for the topic of Chaos Engineering. Authored by A Basiri, N Behnam, R de Rooij, L Hochstein, L Kosewski, J Reynolds, and Casey Rosenthal, this paper provides valuable insights into the subject. [1] It introduces the concept of chaos engineering as a practice for improving the resilience and reliability of complex software systems. The authors provide a comprehensive overview of the principles, methods, and benefits of chaos

engineering. They emphasize the importance of intentionally injecting controlled failures into a system to identify weaknesses and enhance its overall robustness. The authors outline the key components of chaos engineering, including the creation of hypotheses, designing controlled experiments, and the analysis of the system's response to failure scenarios. They discuss various techniques for inducing failures, such as network delays, CPU spikes, and database failures, to stress test the system under realistic conditions. By simulating real-world failures, chaos engineering enables organizations to proactively identify vulnerabilities and address them before they cause significant disruptions. Furthermore, the paper highlights the significance of monitoring and observability in chaos engineering. The authors stress the need for comprehensive monitoring tools and telemetry systems to collect and analyse data during chaos experiments. This data-driven approach helps teams gain insights into system behaviour, identify bottlenecks, and make informed decisions about system improvements.

Moving on, we delve into another paper published in 2020 titled "Identifying and Prioritizing Chaos Experiments by Using Established Risk Analysis Techniques." This insightful work is authored by D Kesim, A van Hoorn, S Frank, and M Häussler. [2] It presents a methodical approach to identify and prioritize chaos experiments, leveraging well-established risk analysis techniques. In the paper "Identifying and Prioritizing Chaos Experiments by Using Established Risk Analysis Techniques," the authors propose a systematic approach to identifying and prioritizing chaos experiments based on established risk analysis techniques. They aim to assist organizations in selecting and conducting chaos experiments that effectively improve the reliability and resilience of their systems while minimizing potential risks. The authors begin by introducing various risk analysis techniques commonly used in other domains, such as failure mode and effects analysis (FMEA), fault tree analysis (FTA), and hazard analysis (HAZOP). They discuss how these techniques can be adapted and applied to chaos engineering, enabling teams to assess the potential impact of chaos experiments on the system and identify the most critical areas for improvement. The paper presents a step-by-step process for conducting risk analysis in the context of chaos engineering. This includes identifying system components, defining failure scenarios, evaluating the severity and likelihood of potential failures, and assigning priorities to chaos experiments based on the identified risks. The authors also emphasize the importance of involving cross-functional teams, including domain experts and stakeholders, to ensure a comprehensive and accurate risk assessment. Additionally, the authors highlight the need for ongoing evaluation and refinement of the risk analysis process. They propose a feedback loop that incorporates lessons learned from previous chaos experiments into subsequent risk assessments, allowing organizations to continuously improve their chaos engineering practices over time.

Lastly, let's explore a research paper titled "Security Chaos Engineering for Cloud Services" by K. A. Torkura, Muhammad I.H. Sukmana, F Cheng, and Christoph Meinel, published in 2018. This study delves into the application of chaos engineering principles to bolster the security of cloud services. [3] The paper explores the application of chaos engineering principles specifically in the context of enhancing the security of cloud services. The authors address the increasing importance of security considerations in cloud environments and propose the integration of chaos engineering as a proactive approach to identify and mitigate security vulnerabilities. The paper starts by discussing the unique challenges faced by cloud service providers regarding security. It highlights the dynamic nature of cloud infrastructures, where constant changes and updates can introduce unintended security weaknesses. The authors argue that traditional security testing methods, such as penetration testing, may not be sufficient to comprehensively assess the security of cloud services, given their complexity and scale. The authors propose the concept of security chaos engineering, which involves intentionally introducing security failures and vulnerabilities into a cloud service environment to test its resilience and identify potential security weaknesses. They discuss various security chaos experiments, such as simulating distributed denial-of-service (DDoS) attacks, injecting malicious code, or tampering with access controls. Furthermore, the paper emphasizes the importance of robust monitoring and incident response capabilities in security chaos engineering. The authors suggest leveraging real-time monitoring tools and security information and event management (SIEM) systems to detect and respond to security incidents effectively. By monitoring the system's response during chaos experiments, organizations can gain valuable insights into their security posture and identify areas for improvement.

III. TERMINOLOGY OVERVIEW

A. Elastic Compute Cloud

Elastic Compute Cloud(EC2) is a key service offered by Amazon Web Services (AWS) that provides resizable compute capacity in the cloud. It allows users to provision virtual servers, known as instances, and enables them to choose the desired configuration, such as the operating system, CPU, memory, and storage capacity. EC2 offers flexibility, scalability, and reliability, allowing businesses to quickly scale their compute resources based on demand. It is commonly used for running applications, hosting websites, and performing various computing tasks in the cloud.

B. Relational Database Service

Relational Database Service(RDS) is a managed database service provided by AWS that simplifies the setup, operation, and scaling of relational databases. It supports popular database engines like Amazon Aurora, MySQL, PostgreSQL, Oracle, and Microsoft SQL Server. RDS automates routine database tasks, such as

backups, software patching, and database maintenance, freeing up users from managing the underlying infrastructure. It offers features like automated backups, high availability, and read replicas, making it easier for businesses to deploy and manage their relational databases in the AWS cloud.

C. *Availability Zone*

An Availability Zone (AZ) is a physically isolated data centre within an AWS region. AWS regions are geographic locations where AWS provides its cloud services. Each region consists of multiple Availability Zones, which are designed to be independent from each other in terms of power, networking, and cooling infrastructure. The purpose of Availability Zones is to provide fault tolerance and high availability for applications and services deployed in the cloud. By distributing resources across multiple Availability Zones, users can protect their applications from failures in a single zone and ensure continuity and resilience.

D. *Fault Injection Simulator*

Fault Injection Simulator (FIS) is an AWS service that allows users to perform controlled experiments to simulate faults and failures in their applications and infrastructure. It helps organizations test the resilience and reliability of their systems by injecting various types of faults, such as network latency, CPU spikes, API errors, or resource exhaustion. With FIS, users can validate how their systems respond to these faults and identify potential weaknesses or bottlenecks. It enables proactive testing and optimization of applications to improve their resilience and readiness for real-world scenarios.

E. *Gremlin*

Gremlin is a chaos engineering tool used to simulate real-world failure scenarios in distributed systems. It helps organizations proactively identify weaknesses by injecting controlled failures such as network outages and latency spikes. With Gremlin, users can assess system resilience, uncover vulnerabilities, and improve overall system robustness and availability. It provides monitoring features to observe system behaviour during chaos experiments, enabling teams to make informed decisions to enhance system performance and reliability.

F. *Boto3*

Boto3 is the official AWS Software Development Kit (SDK) for Python. It provides a simple and intuitive interface to interact with AWS services programmatically. Boto3 allows developers to write Python scripts or applications that can create, configure, and manage AWS resources, such as EC2 instances, S3 buckets, RDS databases, and more. It provides a wide range of APIs and functions to interact with AWS services, making it easier to automate tasks, build applications, and integrate AWS services into Python-based workflows.

G. *Amazon Simple System Manager*

Amazon Systems Manager (SSM) is a fully managed service offered by Amazon Web Services (AWS) that provides a unified interface for managing resources across your AWS environment. SSM simplifies operational tasks, enables automation, and provides a central hub for managing and monitoring infrastructure at scale.

IV. USAGE OF CHAOS ENGINEERING

Chaos engineering is a discipline that involves intentionally injecting controlled disruptions into a system to test its resilience and identify potential weaknesses [5]. When applied to cloud environments, chaos engineering helps organizations proactively uncover and address vulnerabilities, ensuring the robustness and reliability of their cloud-based systems. Here are some key areas where chaos engineering is used in the cloud:

A. *Fault tolerance testing*

Chaos engineering allows organizations to test the fault tolerance of their cloud systems by simulating real-world failure scenarios [6]. By introducing controlled disruptions, such as network outages, instance failures, or service degradation, organizations can observe how their cloud infrastructure and applications respond and recover from these incidents.

B. *Resilience validation*

Chaos engineering helps validate the resilience of cloud systems by deliberately stressing the infrastructure and applications under different load conditions. By introducing heavy traffic, sudden surges, or resource exhaustion, organizations can assess if their cloud environment can handle unexpected spikes in demand while maintaining acceptable performance and availability.

C. *Architecture validation*

Chaos engineering enables organizations to validate the effectiveness of their cloud architecture in handling failures and ensuring high availability [7]. By purposefully inducing failures in specific components or services, organizations can evaluate the impact on the overall system and identify potential single points of failure or architectural weaknesses.

D. *Automated recovery testing*

Chaos engineering can be used to assess the effectiveness of automated recovery mechanisms in cloud environments [4]. By triggering disruptions and failures, organizations can verify if their automated recovery processes, such as auto-scaling, load balancing, or failover mechanisms, are working as intended and can restore the system to a stable state.

E. *Continuous improvement*

Chaos engineering fosters a culture of continuous improvement in cloud environments. By regularly conducting chaos experiments, organizations can gather

valuable insights about their cloud systems, identify weaknesses, and implement necessary improvements to enhance system resilience, reliability, and performance

V. GREMLIN Vs AWS FIS Vs SDK

Gremlin is a chaos engineering platform that allows users to inject controlled failures and disruptions into their systems. It supports various infrastructure components, including cloud services, containers, and hosts.

On the other hand AWS FIS is a service provides by Amazon web Services that enables user to perform controlled fault injection experiments in their AWS infrastructure. FIS integrates with AWS services and provides pre-built templates for conducting chaos experiments. A detailed comparison is done in Table 1 During our study we found that in order to test the system extensively use of backdoor functionality is must. But this can potentially affect the security of our cloud

environment. Imagin a scenario where you are trying to test the system using Gremlin and you create a backdoor for it. If testing is not done in proper secure environment the attacker might use this backdoor to send a malicious payload. Hence use of Backdoor for testing is not a secure approach more over since we are using a third party application to test the application the responsibility of applying security measures lies on user and not the cloud service provider. This increases the workload of the client. Hence use of fault injection simulator is secured way.

But Fault Injection simulator provide limited number of templates to perform the fault injection. For rest of the testing it provide the facility of SSM document for inject the fault. During our study we tried to perform the chaos engineering using the Amazon Software Development Kits (Python SDK). Using it with other monitoring services provided by AWS, we can effectively perform chaos experiments and monitor them.

Feature	Gremlin	Amazon Fault Injection Simulator(FIS)	Software Development Kit(SDK) for chaos Engineering in AWS
Overview	Chaos Engineering as Service platform	Managed Service from AWS	Python-based SDK for automating chaos experiments in AWS
Chaos Engineering Types	Provides various predefined attacks	Offers predefined templates for faults	Allows custom chaos experiment design
Integration with AWS services	Integrates with many AWS services using gremlin agent(Backdoor)	Work seamlessly with AWS ecosystem without backdoor	Provides direct access to AWS APIs and services without backdoor.
Ease of Use	User-friendly web interface	Simple API-driven setup	Requires programming knowledge for custom experiments.
Experiment Customization	Limited customization options	Limited customization options	Highly customizable chaos experiments using Python
Learning Curve	Easy to start with preconfigured attacks	Easy to start with predefined templates	Requires familiarity with Python and AWS services
Managed vs Self-Managed	Managed Service, less control	Managed service form AWS, some control	Self-managed, full control over experiments.
Cost	Subscription-based pricing model	Charged per fault injection	No additional cost for using SDK
Community Support	Strong community support	Backed by AWS support	Community-driven support with Python ecosystem

Table 1 : Comparison Between GREMLIN and AWS FIS and Python SDK for Chaos Engineering

VI. ADVANTAGE AND DISADVANTAGES

Chaos Engineering is widely adopted by companies such as Netflix, Amazon, and Google, offering practical advantages for their cloud applications. When implementing Chaos Engineering using software development kit, several common benefits arise:

- Simplified automation of chaos experiments through a user-friendly interface and comprehensive libraries.
- Seamless integration with various AWS services, providing fine-grained control over infrastructure components.
- Flexibility to design custom chaos experiments tailored to specific AWS services or configurations.
- Access to a rich ecosystem of libraries, frameworks, and community support.
- Alignment with Infrastructure as Code principles for streamlined integration into existing frameworks.

By utilizing the Python SDK for chaos engineering in an AWS environment, teams benefit from simplified automation, seamless integration with AWS services, flexibility for customization, access to a rich ecosystem, and alignment with Infrastructure as Code practices. These advantages empower organizations to efficiently conduct chaos experiments, improve resiliency, and enhance the overall reliability of their AWS-based applications.

While using the Python SDK for Chaos Engineering in an AWS environment offers several advantages, there are also some potential disadvantages to consider:

- Learning Curve: Requires familiarity with Python and AWS services.
- Complexity: Adds complexity to chaos engineering implementation.
- Maintenance and Versioning: Requires ongoing updates and code maintenance.
- Limited Community Support: Specific use case may have fewer resources and support.
- Dependency Management: Managing dependencies and compatibility can be challenging.
- Security and Compliance: Requires attention to security and compliance measures.

VII. CONCLUSION

This paper highlights the considerable benefits of using the Python SDK for chaos engineering in an AWS environment to test the resiliency of AWS services. The Python SDK offers a user-friendly interface for automating chaos experiments, seamless integration with AWS services, customization options for experiments, access to a diverse ecosystem of resources, and alignment with Infrastructure as Code principles. Adopting the

Python SDK empowers organizations to enhance system reliability, optimize service availability, and effectively manage failures. Embracing chaos engineering with the Python SDK in AWS is a highly recommended approach for achieving robust and resilient AWS deployments.

VIII. FUTURE SCOPE

The research on testing the resiliency of AWS services using Chaos Engineering and the Python SDK presents several future opportunities for advancement. Some potential areas of exploration include advanced experiment design methodologies, deeper integration with AWS services, automation and orchestration frameworks, metrics and observability analysis, and the documentation of best practices and case studies. These future directions will contribute to further improving system resilience, enhancing chaos engineering practices, and strengthening the reliability of AWS services.

REFERENCES

- [1] A. Basiri et al., "Chaos Engineering," in *IEEE Software*, vol. 33, no. 3, pp. 35-41, May-June 2016, DOI: 10.1109/MS.2016.60.
- [2] D. Kesim, A. van Hoorn, S. Frank, and M. Häussler, "Identifying and Prioritizing Chaos Experiments by Using Established Risk Analysis Techniques," 2020 IEEE 31st International Symposium on Software Reliability Engineering (ISSRE), 2020, pp. 229-240, DOI: 10.1109/ISSRE5003.2020.00030.
- [3] K. A. Torkura, M. I. H. Sukmana, F. Cheng and C. Meinel, "Security Chaos Engineering for Cloud Services: Work In Progress," 2019 IEEE 18th International Symposium on Network Computing and Applications (NCA), 2019, pp. 1-3, DOI: 10.1109/NCA.2019.8935046.
- [4] C. Nygard, "Release It!: Design and Deploy Production-Ready Software." Pragmatic Bookshelf, 2018. ISBN: 978-1680502398.
- [5] N. Leaver, "Chaos Engineering: Building Confidence in System Behavior through Experiments." O'Reilly Media, 2020. ISBN: 978-1492043864.
- [6] C. Smith and M. Trott, "Learning Chaos Engineering: Explore the world of Chaos Engineering techniques to build resilient and scalable distributed systems." Packt Publishing, 2021. ISBN: 978-1800202390.
- [7] K. Galbraith, "Chaos Engineering: Site reliability through controlled disruption." Manning Publications, 2021. ISBN: 978-1617296252.